

Chapter 18

Machine translation



Machine translation as a task

Machine translation can be formulated as an **optimization problem**

- $w^{(s)}$ is a sentence in a **source language**
- $w^{(t)}$ is a sentence in the **target language**
- Ψ is a **scoring function**.

$$\hat{w}^{(t)} = \operatorname{argmax}_{w^{(t)}} \Psi(w^{(s)}, w^{(t)})$$

This formalism requires **two components**:

1. A **decoding** algorithm for computing $\hat{w}^{(t)}$.
2. A **learning** algorithm for estimating the parameters of the scoring function Ψ

Machine translation : Labeled Data **Problem**

- Labeled translation data usually comes in the form parallel sentences.

$w^{(s)} = A \text{ Vinay le gusta las manzanas.}$
 $w^{(t)} = \text{Vinay likes apples.}$


- A useful feature function would note the translation pairs:
(gusta, likes) (manzanas, apples) (Vinay, Vinay)
- **Problem:** Such word-to-word alignment is **not** given in the data.

Machine translation : Labeled Data **Solutions**

- One **solution** is to treat this alignment as a latent variable; this is the approach taken by classical statistical machine translation (SMT) systems.
- Another **solution** is to model the relationship between $w(t)$ and $w(s)$ through a complex and expressive function; like neural machine translation approach


Vauquois Pyramid

- The **Vauquois Pyramid** is a theory of how translation should be done.
- At the **lowest level**, the translation system operates on **individual words**, but the horizontal distance at this level is large, because languages express ideas differently.
- By moving to **syntactic**, the translation distance is reduced; we then need only produce target-language text from the syntactic representation.



Vauquois Pyramid

- Further up the triangle lies **semantics**; translating between **semantic representations** should be easier still, but mapping between semantics and surface text is a difficult and unsolved problem
- At the **top of the triangle** is **interlingua**, a semantic representation that is so generic that it is identical across all human languages.



Evaluating translations

There are **two main criteria** for a translation:

- **Adequacy**: The translation $w^{(t)}$ should adequately reflect the **linguistic content** of $w^{(s)}$
- **Fluency**: The translation $w^{(t)}$ should read like fluent text in the target language.

Example: Evaluating translations

$w^{(s)} = A \text{ Vinay le gusta Python,}$

$\left. \begin{array}{l} w^{(t)} = \text{Vinay likes Python} \\ w_{\text{gloss}}^{(t)} = \text{To Vinay it like Python} \\ w_3^{(t)} = \text{Vinay debugs memory leaks} \end{array} \right\}$

- **Adequacy:**

- $w_{\text{gloss}}^{(t)}$ or word-for-word translation is adequate because it has all the relevant content.
- The $w_3^{(t)}$ is not adequate base on this criteria.

- **Fluency:** By this criterion, the gloss $w_{\text{gloss}}^{(t)}$ will score poorly, and $w_3^{(t)}$ will be preferred.

Evaluating translations : BLEU Score

- Automated evaluations of machine translations typically merge **Adequacy** and **Fluency** criterias
- Evaluations performed by comparing the system translation with some reference translations, produced by human translators.
- The most popular quantitative metric is **BLEU**, which is based on n-gram precision.
- The n-gram precisions for three hypothesis translations

$$p_n = \frac{\text{number of } n\text{-grams appearing in both reference and hypothesis translations}}{\text{number of } n\text{-grams appearing in the hypothesis translation}}$$

BLEU and Brevity Penalty

- One potential weakness of **BLEU** is that it only measures precision
- The **BLEU** score is then based on the average $\rightarrow \exp \frac{1}{N} \sum_{n=1}^N \log p_n$
- Two modifications are necessary for p_n
 - To avoid computing $\log 0$, all precisions are smoothed to ensure that they are positive
 - Each n-gram in the reference can be used at most once to avoid repetition.
 - Repetition : “to to to to to to” , Reference translation : “to be or not to be”
- **Problem:** precision-based metrics are biased in favor of short translations.
- **Solution :** To avoid this issue, a **brevity penalty “BP”** is applied to translations that are shorter than the reference

BLEU and Brevity Penalty Example

- A reference translation and three system outputs.
- For each output p_n , the precision at each n-gram.
- BP indicates the brevity penalty

	Translation	p_1	p_2	p_3	p_4	BP	BLEU
Reference	<i>Vinay likes programming in Python</i>						
Sys1	<i>To Vinay it like to program Python</i>	$\frac{2}{7}$	0	0	0	1	.21
Sys2	<i>Vinay likes Python</i>	$\frac{3}{3}$	$\frac{1}{2}$	0	0	.51	.33
Sys3	<i>Vinay likes programming in his pajamas</i>	$\frac{4}{6}$	$\frac{3}{5}$	$\frac{2}{4}$	$\frac{1}{3}$	1	.76

Evaluating translations : BLEU Score

- Automated metrics like BLEU have been validated by **correlation** with human judgments of translation quality.
- Nonetheless, it is not difficult to construct examples in which the **BLEU score is high**, yet the **translation is disfluent or carries a completely different meaning from the original**.
- Despite the importance of pronouns for **semantics**, they have a **marginal** impact on BLEU

Pronouns Translation.

- Consider the problem of translating pronouns. Because pronouns refer to specific entities, a single incorrect pronoun can obliterate the **semantics** of the original sentence.
- The problem of **pronoun translation** intersects with issues of **fairness** and **bias**
 - In many languages, the third person singular pronoun is gender neutral.
 - This bias was not directly programmed into the translation model; it arises from statistical tendencies in existing datasets
 - If a dataset has even a slight tendency towards men as doctors, the resulting translation model may produce translations in which doctors are always he, and nurses are always she

Other Translation metrics

- **METEOR** is a weighted F -MEASURE, which is a combination of recall and precision
- **Translation Error Rate (TER)** computes the string edit distance between the reference and the hypothesis
- **RIBES** metric applies rank correlation to measure the similarity in word order between the system and reference translations

Data

- Data-driven approaches to machine translation rely primarily on **parallel corpora**, which are translations at the **sentence level**.
- Many languages have sizable parallel corpora with some high-resource language, but not with each other.
- The high-resource language can then be used as a “pivot” or “bridge”
 - For example, use **Spanish** as a bridge for translation between **Catalan** and **English**.
- For most of the 6000 languages spoken today, the only source of translation data remains the **Judeo-Christian Bible**.
 - While relatively small, at less than a million tokens, the Bible has been translated into more than **2000 languages**

Statistical machine translation : intro

- **Review** : Two main criteria for machine translation **fluency** and **adequacy**
- A **natural modeling** approach is to represent them with **separate scores**

$$\Psi(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) = \Psi_A(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) + \Psi_F(\mathbf{w}^{(t)})$$

- The fluency score Ψ_F need not even consider the source sentence; it only judges $\mathbf{w}^{(t)}$ on whether it is fluent in the target language.
 - This decomposition is advantageous because it makes it possible to estimate the two scoring functions on separate data.
- The fluency model can be estimated from monolingual text in the target language.
 - Large monolingual corpora are now available in many languages, thanks to resources such as Wikipedia.

Noisy channel model

- In the noisy channel model, each **scoring function** is a **log probability**

$$\begin{aligned}\Psi_A(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) &\triangleq \log p_{S|T}(\mathbf{w}^{(s)} | \mathbf{w}^{(t)}) \\ \Psi_F(\mathbf{w}^{(t)}) &\triangleq \log p_T(\mathbf{w}^{(t)}) \\ \Psi(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) &= \log p_{S|T}(\mathbf{w}^{(s)} | \mathbf{w}^{(t)}) + \log p_T(\mathbf{w}^{(t)}) = \log p_{S,T}(\mathbf{w}^{(s)}, \mathbf{w}^{(t)})\end{aligned}$$

- By setting the scoring functions equal to logarithms of the prior and likelihood, their sum is equal to $\log p_{S,T}$, which is **logarithm of joint probability of source and target**.
- The sentence $\hat{\mathbf{w}}^{(t)}$ that maximizes this joint probability is also the maximizer of the conditional probability $p_{T|S}$, making it the most likely target language sentence, conditioned on the source.

Statistical translation modeling : Alignment

- The simplest decomposition of the translation model is **word-to-word**
- Each word in the **source** should be aligned to a word in the **translation**.
- This approach presupposes an **alignment** $\mathcal{A}(w^{(s)}, w^{(t)})$ which contains a **list of pairs** of **source** and **target** tokens.
- Each alignment contains exactly one tuple for each word in the source
- If no appropriate word in the target can be identified for a source word, it is aligned to \emptyset .
- Words in the target can align with multiple words in the source.

Alignment Example

Example:

$$\left\{ \begin{array}{l} \mathbf{w}^{(s)} = A \text{ Vinay } le \text{ gusta } Python \\ \mathbf{w}^{(t)} = \text{Vinay likes } Python \end{array} \right.$$

	<i>A</i>	<i>Vinay</i>	<i>le</i>	<i>gusta</i>	<i>python</i>
<i>Vinay</i>					
<i>likes</i>					
<i>python</i>					

$$\left\{ \begin{array}{l} \mathcal{A}(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) = \{(A, \emptyset), (Vinay, Vinay), (le, likes), (gusta, likes), (Python, Python)\} \\ \mathcal{A}(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) = \{(A, Vinay), (Vinay, likes), (le, Python), (gusta, \emptyset), (Python, \emptyset)\} \end{array} \right.$$

Alignment and Translation

- The **joint probability** of the **alignment** and the **translation**:

$$\begin{aligned} p(\mathbf{w}^{(s)}, \mathcal{A} \mid \mathbf{w}^{(t)}) &= \prod_{m=1}^{M^{(s)}} p(w_m^{(s)}, a_m \mid w_{a_m}^{(t)}, m, M^{(s)}, M^{(t)}) \\ &= \prod_{m=1}^{M^{(s)}} p(a_m \mid m, M^{(s)}, M^{(t)}) \times p(w_m^{(s)} \mid w_{a_m}^{(t)}) \end{aligned}$$

Statistical translation modeling

This probability model makes two key assumptions:

- The alignment probability factors across tokens
- This means that each alignment decision is independent of the others, and depends only on the index m , and the sentence lengths $M^{(s)}$ and $M^{(t)}$

$$\begin{aligned} p(\mathbf{w}^{(s)}, \mathcal{A} | \mathbf{w}^{(t)}) &= \prod_{m=1}^{M^{(s)}} p(w_m^{(s)}, a_m | w_{a_m}^{(t)}, m, M^{(s)}, M^{(t)}) \\ &= \prod_{m=1}^{M^{(s)}} p(a_m | m, M^{(s)}, M^{(t)}) \times p(w_m^{(s)} | w_{a_m}^{(t)}) \end{aligned}$$

$$p(\mathcal{A} | \mathbf{w}^{(s)}, \mathbf{w}^{(t)}) = \prod_{m=1}^{M^{(s)}} p(a_m | m, M^{(s)}, M^{(t)})$$

- The translation probability also factors across tokens
- so that each word in $\mathbf{w}^{(s)}$ depends only on its aligned word in $\mathbf{w}^{(t)}$. This means that translation is word-to-word, ignoring context.

$$p(\mathbf{w}^{(s)} | \mathbf{w}^{(t)}, \mathcal{A}) = \prod_{m=1}^{M^{(s)}} p(w_m^{(s)} | w_{a_m}^{(t)})$$

Statistical translation modeling

- To translate with such a model, we could sum or max over all possible alignments
- The term $p(\mathcal{A})$ defines the prior probability over alignments.

$$\begin{aligned} p(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) &= \sum_{\mathcal{A}} p(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}, \mathcal{A}) \\ &= p(\mathbf{w}^{(t)}) \sum_{\mathcal{A}} p(\mathcal{A}) \times p(\mathbf{w}^{(s)} \mid \mathbf{w}^{(t)}, \mathcal{A}) \\ &\geq p(\mathbf{w}^{(t)}) \max_{\mathcal{A}} p(\mathcal{A}) \times p(\mathbf{w}^{(s)} \mid \mathbf{w}^{(t)}, \mathcal{A}) \end{aligned}$$

Neural machine translation

- Neural network models for machine translation are based on the encoder-decoder architecture
- The **encoder** network converts the source language sentence into a vector or matrix representation
- The **decoder** network then converts the encoding into a sentence in the target language

$$\begin{aligned} z &= \text{ENCODE}(w^{(s)}) \\ w^{(t)} \mid w^{(s)} &\sim \text{DECODE}(z), \end{aligned}$$

Where the second line means that the function $\text{DECODE}(z)$ defines the conditional probability $p(w^{(t)} \mid w^{(s)})$

Neural machine translation

- The **decoder** is typically a recurrent neural network, which generates the **target** language sentence one word at a time, while recurrently updating a hidden state.
- The encoder and decoder networks are trained **end-to-end** from **parallel sentences**.
- If the **output** layer of the decoder is a **logistic** function, then the entire architecture can be trained to maximize the conditional log-likelihood:

$$\log p(\mathbf{w}^{(t)} \mid \mathbf{w}^{(s)}) = \sum_{m=1}^{M^{(t)}} \log p(w_m^{(t)} \mid \mathbf{w}_{1:m-1}^{(t)}, \mathbf{z})$$
$$p(w_m^{(t)} \mid \mathbf{w}_{1:m-1}^{(t)}, \mathbf{w}^{(s)}) \propto \exp \left(\beta_{w_m^{(t)}} \cdot \mathbf{h}_{m-1}^{(t)} \right)$$

Neural machine translation

- When the **output** layer of decoder is **logistic** function:

$$\log p(\mathbf{w}^{(t)} \mid \mathbf{w}^{(s)}) = \sum_{m=1}^{M^{(t)}} p(w_m^{(t)} \mid \mathbf{w}_{1:m-1}^{(t)}, \mathbf{z})$$
$$p(w_m^{(t)} \mid \mathbf{w}_{1:m-1}^{(t)}, \mathbf{w}^{(s)}) \propto \exp \left(\beta_{w_m^{(t)}} \cdot \mathbf{h}_{m-1}^{(t)} \right)$$

- Where the hidden state $\mathbf{h}_{m-1}^{(t)}$ is a **recurrent function** of the previously generated text $\mathbf{w}_{1:m-1}^{(t)}$ and the encoding \mathbf{z} .

$$p(w_m^{(t)} \mid \mathbf{w}_{1:m-1}^{(t)}, \mathbf{w}^{(s)}) \propto \exp \left(\beta_{w_m^{(t)}} \cdot \mathbf{h}_{m-1}^{(t)} \right)$$



$$w_m^{(t)} \mid \mathbf{w}_{1:m-1}^{(t)}, \mathbf{w}^{(s)} \sim \text{SoftMax} \left(\beta \cdot \mathbf{h}_{m-1}^{(t)} \right)$$

- Where $\beta \in \mathbb{R}^{(V^{(t)} \times K)}$ is the **matrix of output word vectors** for the $V^{(t)}$ words in the target language vocabulary.

Sequence-to-Sequence Model

- The **simplest** encoder-decoder architecture is the sequence-to-sequence model
- Sequence-to-sequence translation is nothing more than wiring together two LSTMs: one to **read the source**, and another to **generate the target**.
- In this model, the encoder is set to the final hidden state of a long short-term memory (LSTM) on the source sentence

$$\begin{aligned} h_m^{(s)} &= \text{LSTM}(x_m^{(s)}, h_{m-1}^{(s)}) \\ z &\triangleq h_{M^{(s)}}^{(s)}, \end{aligned}$$

- Where $x_m^{(s)}$ is the embedding of source language word $w_m^{(s)}$.
- The encoding then provides the initial hidden state for the decoder LSTM

$$\begin{aligned} h_0^{(t)} &= z \\ h_m^{(t)} &= \text{LSTM}(x_m^{(t)}, h_{m-1}^{(t)}), \end{aligned}$$


Tweaks on Sequence-to-Sequence Model

- The model works better if the source sentence is **reversed**, reading from the end of the sentence back to the beginning.
- In this way, the words at the **beginning** of the **source** have the **greatest** impact on the **encoding z** , and therefore impact the words at the **beginning** of the **target** sentence.
- Advance encoding models, such as **neural attention** has **eliminated** the need for **reversing** the source sentence.

Tweaks on Sequence-to-Sequence Model : pt2

- The encoder and decoder can be implemented as **deep LSTMs**, with multiple layers of hidden states
- Each hidden state $h_m^{(s,i)}$ at layer i is treated as input to an LSTM at layer $i + 1$

$$h_m^{(s,1)} = \text{LSTM}(x_m^{(s)}, h_{m-1}^{(s)})$$
$$h_m^{(s,i+1)} = \text{LSTM}(h_m^{(s,i)}, h_{m-1}^{(s,i+1)}), \quad \forall i \geq 1$$



Tweaks on Sequence-to-Sequence Model : pt3

- Significant improvements can be obtained by creating an ensemble of translation models, each trained from a different random initialization.
- For an ensemble of size **N**, the per-token decoding probability is set equal to:

$$p(w^{(t)} \mid z, w_{1:m-1}^{(t)}) = \frac{1}{N} \sum_{i=1}^N p_i(w^{(t)} \mid z, w_{1:m-1}^{(t)})$$


- where p_i is the decoding probability for model i

Model comparison

- Statistical Translation approach is **compositional**
- Encoder-decoder models are **contextualized**
- **Question:** Is it possible for translation to be both **contextualized** and **compositional**?
- **Answer:** Yes, One approach is to augment neural translation with an **attention mechanism**.

Neural Attention


- **Neural attention** makes it possible to integrate **alignment** into the **encoder-decoder architecture**.
- In general, attention can be thought of as using a query to select from a memory of **key-value pairs**.
- However, the query, keys, and values are all vectors, and the entire operation is differentiable
- A general view of neural attention:




Neural Attention

- For each key n in the memory, score $\psi_{\alpha}(m, n)$ gets computed with respect to the query m
- $\psi_{\alpha}(m, n)$ score is a function of the **compatibility** of **key** and **query**
- $\psi_{\alpha}(m, n)$ can be computed using a **small feedforward neural network**.
- The **vector of scores** is passed through an **activation function**, like SoftMax.
- Output of activation function is a vector of non-negative numbers:

$$[\alpha_{m \rightarrow 1}, \alpha_{m \rightarrow 2}, \dots, \alpha_{m \rightarrow N}]^T$$



Neural Attention



- Output of activation function is a vector of non-negative numbers $[\alpha_{m \rightarrow 1}, \alpha_{m \rightarrow 2}, \dots, \alpha_{m \rightarrow N}]^T$
- $[\alpha_{m \rightarrow 1}, \alpha_{m \rightarrow 2}, \dots, \alpha_{m \rightarrow N}]^T$ has length of N, equal to the size of the memory.
- Each value in the memory \mathbf{v}_n is multiplied by the attention $\alpha_{m \rightarrow n}$; the **sum** of these scaled values is the output
- The dotted box indicates that each $\alpha_{m \rightarrow n}$ can be viewed as a **gate** on value n
- **Extreme case:** $\alpha_{m \rightarrow n} = 1$ and $\alpha_{m \rightarrow n'} = 0$ for all other n' .
Then the attention mechanism selects the value \mathbf{v}_n from memory.

Neural Attention : Decoding

- At each step m in decoding, attentional state is computed by executing a query, which is equal to the **state of decoder**, $\mathbf{h}_m^{(t)}$. The resulting compatibility scores :

$$\psi_{\alpha}(m, n) = \mathbf{v}_{\alpha} \cdot \tanh(\Theta_{\alpha}[\mathbf{h}_m^{(t)}; \mathbf{h}_n^{(s)}])$$

- The function ψ is thus a **two-layer feedforward neural network**, with weights \mathbf{v}_{α} on the **output layer**, and weights Θ_{α} on the **input layer**.
- To convert these **scores** into **attention weights**, we apply an **activation function**, which can be **vector-wise SoftMax** or an **element-wise sigmoid**

Neural Attention : Activation functions

- Vector-wise SoftMax attention:

$$\alpha_{m \rightarrow n} = \frac{\exp \psi_{\alpha}(m, n)}{\sum_{n'=1}^{M^{(s)}} \exp \psi_{\alpha}(m, n')}$$

- Element-wise sigmoid attention:

$$\alpha_{m \rightarrow n} = \sigma(\psi_{\alpha}(m, n))$$

$$\mathbf{c}_m = \sum_{n=1}^{M^{(s)}} \alpha_{m \rightarrow n} \mathbf{z}_n$$

$$\tilde{\mathbf{h}}_m^{(t)} = \tanh(\Theta_c[\mathbf{h}_m^{(t)}; \mathbf{c}_m])$$

$$\mathbf{p}(w_{m+1}^{(t)} | \mathbf{w}_{1:m}^{(t)}, \mathbf{w}^{(s)}) \propto \exp\left(\beta_{w_{m+1}^{(t)}} \cdot \tilde{\mathbf{h}}_m^{(t)}\right)$$

- Attention α is then used to compute a context vector \mathbf{c}_m by taking a weighted average over the columns of \mathbf{Z} .
- $\alpha_{m \rightarrow n} \in [0, 1]$ is the **amount of attention** from **word** m of the **target** to **word** n of the **source**.

The decoder state $\mathbf{h}_m^{(t)}$ is concatenated with the context vector, forming the input to compute a final output $\tilde{\mathbf{h}}_m^{(t)}$.